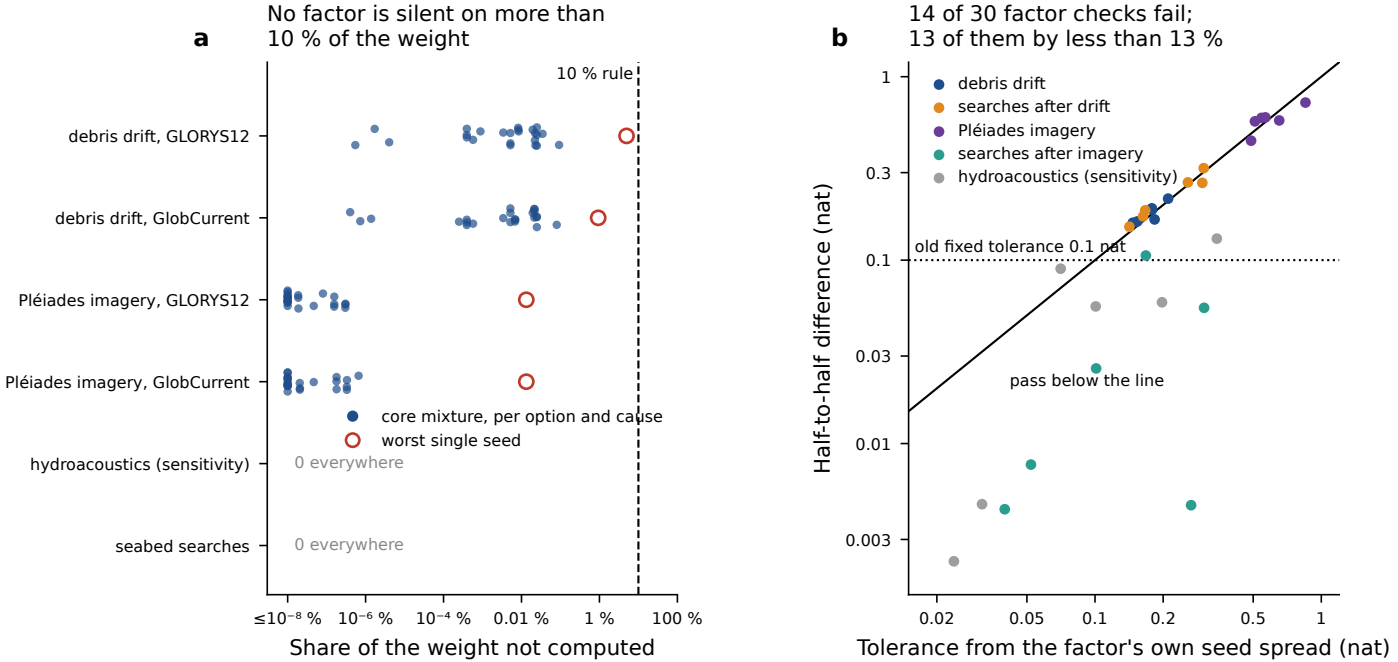

(i) Panel a: for each data source, the part of the probability on which it gave no value. The rule puts a source with more than 10 % in a separate column. No source is near that. Panel b: each dot is one source and one 00:19 data option. A dot above the line means the two halves of the seeds disagree by more than the seeds' own scatter allows. Core run I is not converged; the end-of-flight uncontrolled arm is not calibrated; equal family prior.

(ii) Core run I, 4 families x 8 seeds x 3.2 M impacts (102.4 M rows), equal-prior core mixture (ruling C), cause model q = 1/3 with C2 silent in (b), all cause models in (a); not-computed share = pre-composition (base 00:19-cell) weight of NaN rows, seeds pooled by base evidence; rows carried at the (seed, stratum) mean likelihood ratio. (b) statistic |Delta(1-4) - Delta(5-8)| of the evidence-pooled log-evidence increment (compose crate convention); tolerance = 2 sqrt(1/4 + 1/4) x SD over 8 replicates of the per-replicate increment. Ocean models mixed 0.5(L_G + L_C); Pléiades rho 0.25, conditional on H, A_scene constant omitted; searches rho 0.05 (Ocean Infinity excluded); hydro lnL_soft labelled sensitivity. Australian non-recovery not composed. Script report_r15.py (pass 2b). Imagery shares below 10⁻⁸ % are drawn at the left edge.